




Applied Machine Learning

The Data Science Lifecycle

Main phases



Learning goals

- Know phases of data science lifecycle
- Identify essential concepts, tools, and methods for each phase

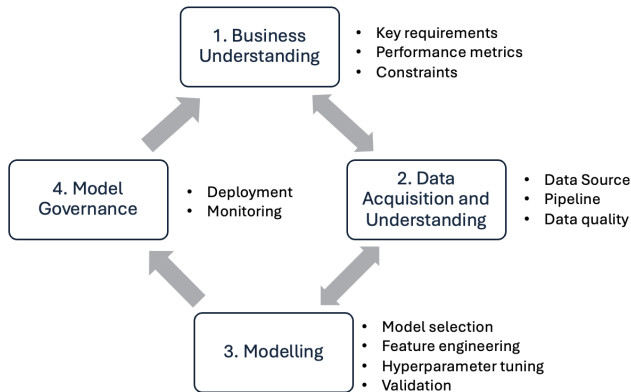
THE DATA SCIENCE LIFECYCLE

► “CRISP-DM” n.d.



Lifecycle Management: Process of managing a (product) lifecycle from inception, through engineering, design and manufacturing to deployment and eventual disposal.

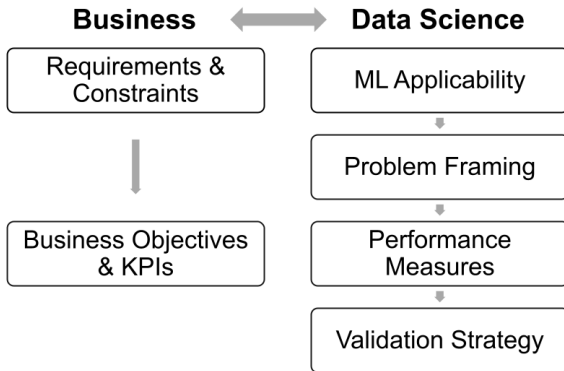
CRISP-DM (Cross-Industry Standard Process for Data Mining): A widely used framework that formalizes the iterative process of turning data into knowledge.





1. Business Understanding

OVERVIEW BUSINESS UNDERSTANDING



REQUIREMENTS & CONSTRAINTS

Key question: What is the end goal of the project?

- Define the intended **output**:
 - Report, dashboard, or interactive tool?
 - Scalable pipeline for production use?



REQUIREMENTS & CONSTRAINTS



Key question: What is the end goal of the project?

- Define the intended **output**:
 - Report, dashboard, or interactive tool?
 - Scalable pipeline for production use?
- Identify constraints
 - Ethical: Avoid sensitive data use (e.g., address may proxy ethnicity)
 - Legal: Exclude protected attributes (e.g., gender in automated hiring)
 - Reproducibility: Ensure the analysis is fully reproducible in a single step
 - Explainability: All model decisions must be explainable to the user
 - Prediction Latency: Real-time predictions required
 - Model Size: Stay within specified memory constraints

REQUIREMENTS & CONSTRAINTS

1. ML Applicability: Is ML justified?

- Explore non-ML baselines first (rules, heuristics, simple analytics)
- Use ML only if it adds clear value over simpler/existing alternatives



REQUIREMENTS & CONSTRAINTS




1. ML Applicability: Is ML justified?

- Explore non-ML baselines first (rules, heuristics, simple analytics)
- Use ML only if it adds clear value over simpler/existing alternatives

2. Problem Framing: How to frame the task?


- *Imbalanced data*: classification vs. anomaly detection?
- *Pred. maintenance*: failure probability (classif.) vs. remaining lifetime (regr.)?
- Select the ML task that aligns with your KPI and data characteristics



BUSINESS OBJECTIVES: PERFORMANCE MEASURES

- How to translate key performance indicators (KPIs) to metrics?

- **Spam detection:** Job application flagged as spam = unacceptable
⇒ Occasional inbox spam is fine ⇒ minimize FP ⇒ maximize **precision**
- **Airport security:** Weapon flagged as laptop = unacceptable
⇒ Occasional false alarm is fine ⇒ minimize FN ⇒ maximize **recall**



BUSINESS OBJECTIVES: PERFORMANCE MEASURES



- **How to translate key performance indicators (KPIs) to metrics?**
 - **Spam detection:** Job application flagged as spam = unacceptable
⇒ Occasional inbox spam is fine ⇒ minimize FP ⇒ maximize **precision**
 - **Airport security:** Weapon flagged as laptop = unacceptable
⇒ Occasional false alarm is fine ⇒ minimize FN ⇒ maximize **recall**
- **Multiple objectives:** What if improving one metric worsens another?
 - *Scalarization:* Combine into one score (e.g., F_1)
 - *Pareto/frontier:* Analyze trade-off without collapsing (e.g., ROC curve)

True Values	Predicted Values		
	Pos	Neg	
Pos	TP	FN	$\text{Sensitivity/Recall/TPR} = \frac{TP}{TP + FN}$ $\text{Specificity} = \frac{TN}{TN + FP}$
Neg	FP	TN	
	$\text{Precision} = \frac{TP}{TP + FP}$	$\text{NPV} = \frac{TN}{TN + FN}$	$F_1 = \frac{2 \text{ Precision Recall}}{\text{Precision} + \text{Recall}}$

BUSINESS OBJECTIVES: VALIDATION STRATEGY



Key question: How to validate the model reliably and realistically?


1. Strategic Considerations

- **What should your model generalize to?**

- *Blocking factors* (e.g., customers, hospitals, product types)
- *Temporal effects* (e.g., seasonality, trends)
- *Spatial effects* (e.g., regions, branches)

- **Simulate realistic deployment:**

- **Time series:** Train on data from 2016–2020, deploy in 2021
- **Validation:** Mimic deployment – train on 2016–2019, validate on 2020



BUSINESS OBJECTIVES: VALIDATION STRATEGY



Key question: How to validate the model reliably and realistically?

1. Strategic Considerations

- **What should your model generalize to?**

- *Blocking factors* (e.g., customers, hospitals, product types)
- *Temporal effects* (e.g., seasonality, trends)
- *Spatial effects* (e.g., regions, branches)

- **Simulate realistic deployment:**

- **Time series:** Train on data from 2016–2020, deploy in 2021
- **Validation:** Mimic deployment – train on 2016–2019, validate on 2020

2. Beware of Data Leakage

- **Target Leakage:** Features leak future info not available at prediction time
 - Example: "Total Sales in Next Month" used to predict this month
- **Train-Test Leakage:** Information from the test set is used during training
 - Example: Scaling the full dataset before splitting



2. Data Acquisition and Understanding

DATA SOURCES AND DATA VERSIONING

Key question: What data is available and where does it come from?



- **Internal vs. External:** External feeds may change format or be discontinued
- **Structured vs. Unstructured:** Tables vs. free-text, images, audio, etc.
- **Databases vs. Files:** DBs offer querying, consistency and security
- **On-Premise vs. Cloud:** Governed by regulations, cost and team expertise

Data Versioning:

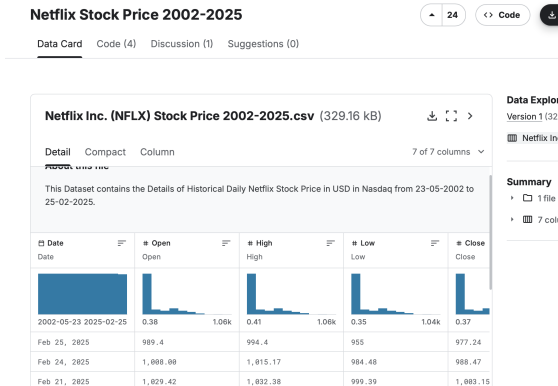
- Data can evolve - use version control for data, like `git` for code
- Tools: DVC (Data Version Control), Pachyderm





Datasheets summarize key metadata about a dataset's content, structure, provenance, and known limitations to support informed and responsible use.

- Why and how was the data collected?
- What are its limitations and ethical concerns?
- For which tasks is it suitable or unsuitable?





Source: ► "Kaggle Datasets" 2025

DATA EXPLORATION



Key question: What features are in the data, and what do they reveal?

- **Build a data dictionary:**
 - Feature meaning, units, valid ranges/category-levels
 - Missing value encoding and special codes
- **Examine structure and quality to gain understanding:**
 - Compute (stratified) summary statistics and inspect missingness
 - Plot univariate and bivariate distributions
 - Identify outliers, anomalies, imbalances, redundancy, or high correlation
- **Watch for confounders:** unmeasured variables may distort associations



DATA QUALITY

Key question: Is the data fit for modeling? *Garbage in, garbage out.*

- Assess missingness: How many values are missing, and where?
- Detect implausible or inconsistent values (requires domain expertise)
- Standardize formats (e.g., date/time parsing)
- Drop near-constant or non-informative features
- Remove redundancy and obsolete variables

Watch for data bias:


- *Sampling bias:* E.g., a survey at 10am in the city center misses working adults
- *Response bias:* Online reviews often reflect extreme opinions, not the median



3. Modeling



FEATURE ENGINEERING



Feature Engineering transforms raw data into informative input variables for modeling.


- **Transformation:** Scaling, log-transform, one-hot encoding
- **Selection:** Wrapper and filter methods, model-based importance
- **Dimensionality Reduction:** PCA, t-SNE
- **Creation:** Aggregation, binning, interaction terms

Best Practices:

- Avoid overengineering: irrelevant features introduce noise
- Fit transformations **only on training data**; reuse parameters for test data
- Prevent leakage: never use test data in feature construction



MODEL TRAINING



Treat model training like a scientific experiment:


- **Track all inputs precisely:**

- Dataset version
- Algorithm and hyperparameters
- Code version (e.g., Git commit/tag)

- **Proceed iteratively:**

- *Start simple* – build a fast, interpretable baseline
- Gradually increase model complexity and validate each improvement


MODEL CONFIGURATION



- How do you track the exact configuration of a model experiment?
- How do you record small changes between experiments?
- Store in plain text (e.g., `.yaml`, `.json`) for readability and version control
- Recommended tool: structured configuration managers
 - **Python:** ▶ "Hydra" n.d. - structured, hierarchical configuration management
 - **R:** ▶ "R config package" n.d. - environment-based YAML setup




MODEL EVALUATION



- Select models based on predefined **metrics** (e.g., accuracy, AUC, RMSE)
- Use the **validation or test set** in line with your validation strategy
- Assess **robustness**: How does the model respond to noise, perturbations, or adversarial inputs?
 - ⇒ Add synthetic noise and measure changes in model parameters/performance



INTERPRETATION AND FAIRNESS



To understand and trust your model, apply methods from **Interpretable Machine Learning (IML)**:

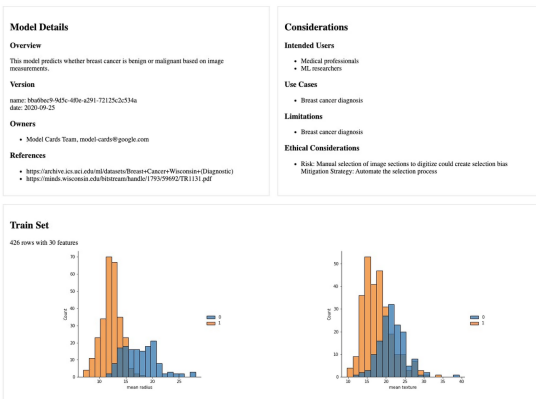
- **Feature Importance**: Which features drive predictions globally?
- **Effects & Interactions**: Visualize marginal effects and feature interactions
- **Local Explanations**: Explain individual predictions
- **Bias Detection**: Reveal fairness issues and unintended discrimination
 - Example: Amazon's 2018 hiring tool penalized resumes with "women's" due to biased training data ▶ "The Guardian" 2018



Model Card: Documentation of the model similar to a Datasheet for data

- **Model Details:** Algorithm, parameters, model version
- **Intended Use:** Intended use and out-of-scope use cases
- **Performance:** Key metrics, validation approaches
- **Analyses:** Feature/target distributions in train/test data
- **Ethical Considerations:** Failure modes, fairness concerns, inappropriate use

Model Card for Breast Cancer Wisconsin (Diagnostic) Dataset






4. Model Governance

MODEL GOVERNANCE

- Model Governance includes **people, processes, and technologies** needed to manage and protect your ML models.
- Ensures long-term maintainability and adaptation as requirements evolve.
- Two extremes must be avoided:
 - **Repression**: too many rules, rigid standards, excessive control
 - **Chaos**: too much speed, freedom, creativity, and unstructured change

Model governance steps are:



"Developing and deploying ML systems is relatively fast and cheap, but maintaining them over time is difficult and expensive."

► "Hidden Technical Debt in ML Systems" n.d.



DEPLOYMENT

Deployment means making a model accessible to users (humans or other systems).

- Common deployment strategies:
 - **One-off**: Manually trained and deployed once; no regular updates
 - **Batch**: Periodically retrained on recent data (e.g., weekly updates)
 - **Real-time**: Continuously updated on streaming data; often latency-critical
- For frequent updates, establish model **versioning and lifecycle tracking**
 - Tools like ["MLflow" n.d.](#) provide version control and a model registry



MONITORING

Monitoring ensures that deployed models work reliably and flags issues in production:

- **Pipeline Failures:** Bugs in data ingestion or transformation logic
- **Data Drift:** Changes in input distribution
 - Example: An image classifier trained on summer data fails when images contain winter snow
 - Also affects evolving targets (e.g., new categories or value ranges)
 - **Remedy:** Label sufficient new data and retrain the model
- **Concept Drift:** Changes in the meaning of labels
 - Example: Redefining “high blood pressure” from 140mmHg to 135mmHg
 - **Remedy:** Relabel historical data and retrain



USER FEEDBACK, REINFORCEMENT



Handover: At project completion, engage users and transfer responsibilities:

- **System Validation:** Confirm with stakeholders that the solution meet their needs
- **Project Handover:** Transfer system ownership to the production team

Feedback loops occur when model outputs influence future inputs:

- **Direct Loops:** A model affects the data it later learns from
 - Example: Netflix recommends certain movies, then only learns from user responses to those
- **Hidden Loops:** Multiple models indirectly affect each other via the environment
 - Example: Flash crashes triggered by interacting trading algorithms amplifying market reactions