

# Applied Machine Learning

## Parallelization: Batchtools package



### Learning goals

- Understand parallelization concepts
- Introduction to batchtools package

# PARALLELIZATION



- **Goal:** Minimize computation time by distributing tasks across CPUs/GPUs.  
⇒ Speedup is ideally linear but often limited by overhead and dependencies.
- Debugging parallel code is especially hard.
- Coding discipline is even more important to minimize errors and frustration.
- **What can be easily parallelized?**
  - Independent replications
  - Resampling, cross-validation
  - Model averaging
  - Parameter variations in simulations ...
  - "Single program, multiple data"
  - Everything expressible as a loop of independent iterations  
(if you can write it with `(1 | m | ) apply`, you are fine)
- **Many statistical problems are "embarrassingly parallel"**

# NAIVE BATCH COMPUTING (NON-CLUSTER)




## Workflow on multicore machines:

- Write standalone script(s) to run jobs and save outputs.
- Hard-code parameters or pass via Command-Line Interface (CLI).
- Log in via SSH; run with `R CMD BATCH myscript.R`.
- Use `nohup`, `screen`, or `tmux` to persist after logout.
- Manually start jobs when CPUs free up.
- Check completion and errors by hand.
- Write scripts to merge results.

## Drawbacks:

- No resource management, automation, or fair scheduling.
- Poor scalability; hard to debug parallel issues.
- No guarantees for reproducibility (e.g., seeding).

# HIGH-PERFORMANCE COMPUTING (HPC) CLUSTERS



[www.oxygen-icons.org](http://www.oxygen-icons.org)

- Users access a gateway server (head node).
- Cluster = multiple nodes managed by a scheduler (e.g., SLURM).
- Scheduler assigns jobs to nodes via a queuing system.
- Nodes share a common file system.

# MANUAL WORKFLOW ON A HPC CLUSTER



## Resource Specification:

- Define CPU count, memory, runtime, partition.
- Set command (e.g., `R CMD BATCH script.R`).

## Manual Tasks:

- Submit jobs via CLI or shell scripts.
- Monitor with tools like `squeue`.
- Write aggregation scripts for results.

## Typical Workflow:

- Unroll R loops into single-iteration scripts.
- Auto-generate job and script files per task.
- Submit jobs; crawl logs and outputs.

## Handling Issues:

- Kill + resubmit on failure.
- Adjust resources on wall-time hits.
- Full rerun for changes in data or params.

# BATCHTOOLS OVERVIEW



- R package for structured access to batch systems.
- Built around Map-Reduce: apply algorithms to many problems.
- Full control from R: submit, monitor, kill jobs.
- Persistent state: resume and audit large experiments.
- Convenient debugging and result collection.
- Supports reproducibility across hardware and job schedulers.
- Supports multiple execution backends:
  - **Interactive:** Run jobs directly in the current R session
  - **Multicore:** Parallel execution on local CPU cores
  - **SSH:** Offload jobs to remote machines via SSH
  - **HPC schedulers:** SLURM, Torque/PBS, Load Sharing Facility (LSF), etc.

**Project Page:** <https://github.com/mllg/batchtools>

**Paper:** <https://doi.org/10.21105/joss.00135>

# CREATING AND CONFIGURING A REGISTRY



## Purpose:

- A registry object is used to access and exchange information: file paths, job parameters, and computational events, ...
- Stores all data in a single, portable directory for easy tracking and reproducibility.

## Initialization of a new registry:

```
library(batchtools)
reg = makeRegistry(
  file.dir = "registry", # Directory accessible on all nodes
  seed = 1               # Initial seed for reproducibility
)
```

## Configure the system:

```
# Set interactive mode and start jobs in external R sessions
reg$cluster.functions = makeClusterFunctionsInteractive(external = TRUE)
```

- Each supported system has its own makeClusterFunctions\* function.

## Load an existing registry to continue work:

```
loadRegistry("registry")
```

# DEFINE JOBS

batchMap:

- Like `lapply` or `mapply`
- $(x_1, x_2) \times (y_1, y_2) \rightarrow (f(x_1, y_1), f(x_2, y_2))$
- 10 jobs to calculate  $1 + 9, 2 + 8, \dots, 9 + 1$

```
map = function(i, j) i + j  
ids = batchMap(fun = map, i = 1:9, j = 9:1, reg = reg)
```

- Stores function on file system
- Creates jobs as rows in a `data.table`
- Parameters also serialized into the `data.table` for fast access
- All jobs get unique positive integers as IDs
- `reg` = can be omitted in most cases. See `?getDefaultRegistry`.





# SUBSET JOBS



## Query Job IDs by Status and Parameters

- Use `find*` functions to query job IDs by computational status:
  - `findError` to get job IDs for failed jobs
  - `findDone` to get job IDs for successful jobs
  - `findNotSubmitted` to get job IDs in order resume jobs
  - ...
- Query job IDs by parameters with `findJobs(pars)` (here: `i` and `j`), e.g.,:

```
job = findJobs(i == 2 & j == 8)
job
## Key: <job.id>
##   job.id
##   <int>
## 1:      2
```

- Pass the `data.table` containing the `job.ids` to functions interacting with the batch system, e.g., `submitJobs(ids = job)`

# SUBMIT JOBS



- Creates R script files and job description files on the fly
- Resources can be provided as named list

```
# 1 hour maximal execution time, about 2 GB of RAM
res = list(walltime = 60*60, memory = 2000)

# ... and submit
submitJobs(resources = res)
```

- Submits all jobs per default
- Subsets of jobs can be providing as `data.table` or vector

```
submitJobs(ids = 1:5, resources = res)
```

- Collect/reduce results:

```
# get results of each job in a list
reduceResultsList(ids = findDone())

# get result of single job
loadResult(id = 1)
```

# SUPERVISE AND DEBUG



- Quick overview of what is going on: `getStatus()`

```
## Status for 9 jobs at 2019-10-10 17:49:48:  
## Submitted      : 9 (100.0%)  
## -- Queued       : 0 (  0.0%)  
## -- Started      : 9 (100.0%)  
## ---- Running    : 0 (  0.0%)  
## ---- Done       : 9 (100.0%)  
## ---- Error      : 0 (  0.0%)  
## ---- Expired    : 0 (  0.0%)
```


- Display log files with a customizable pager (`less`, `vi`, ...):  
`showLog(findErrors()[1])`
- You can also `grepLogs(pattern)`
- Found a bug? `killJobs(findRunning())`
- Run a job in the current R session: `testJob(id)`

# EXPERIMENTS IN BATCHTOOLS



- **Purpose:** Abstraction for typical statistical tasks.
- **Applying Algorithms to Problems:**
  - Ideal for simulations, benchmark experiments, sensitivity analyses,  
...
  - Simplifies workflow with a focus on job definition.
- **Scenarios:**
  - Compare machine learning algorithms on multiple datasets.
  - Compare one/many estimation procedure(s) on simulated data.
  - Compare optimizers on various objective functions.

# ABSTRACTION OF COMPUTER EXPERIMENTS



- **Problem Definition:**

- **Static part:** Immutable R objects (e.g., matrices, data frames).
- **Dynamic part:** Arbitrary R functions (e.g., transformations of static objects, data extraction from external sources, data generation functions).

- **Parametrization:** Specify experimental designs for problems and algorithms.

- **Seeding and Reproducibility:**

- Each step is automatically seeded.
- Random seeds are stored in a database for reproducibility.

# EXPERIMENT DEFINITION STEPS

- Add problems to registry: `addProblem`
  - Efficient storage: Separation of static (data) and dynamic (instance) problem parts.
- Add algorithms to registry: `addAlgorithm`
  - Problem instance gets passed to algorithm
  - Can be connected with an experimental design (function parameters)
  - Return value will be saved on the file system
- Add experiments to registry: `addExperiments`
  - Experiment: problem instance + algorithm + algorithm parameters
  - Job: Experiment + replication number



# A SIMPLE EXAMPLE



```
reg = makeExperimentRegistry("test_reg")
addProblem(name = "p1", data = 1, seed = 1,
  fun = function(data, job) runif(data))
addAlgorithm(name = "a1",
  fun = function(job, data, instance) 2 * instance)
addAlgorithm(name = "a2",
  fun = function(job, data, instance) data + instance)
addExperiments(repls = 2)
submitJobs()
res = reduceResultsDataTable()
getJobPars()[res]
```

## Key: <job.id>

##	job.id	problem	prob.pars	algorithm	algo.pars	result
##	<int>	<char>	<list>	<char>	<list>	<list>
## 1:	1	p1	<list[0]>	a1	<list[0]>	0.5310
## 2:	2	p1	<list[0]>	a1	<list[0]>	0.3698
## 3:	3	p1	<list[0]>	a2	<list[0]>	1.2655
## 4:	4	p1	<list[0]>	a2	<list[0]>	1.1849

# SUMMARY



- **Reproducibility:**

- Every computation is seeded.
- Seeds are stored in a `data.table`.

- **Extensibility:**

- Easily add more problems or algorithms.
- Try different parameters or increase replications at any stage.

- **Portability:** Data, algorithms, results, and job information in a single directory.

- **Exchangeability:** Share your file directory to allow others to extend your study with their data sets and algorithms.

- Simplifies working with batch systems.
- Control batch systems interactively from within R (no shell required).
- Facilitates reproducible research.
- Enables easy exchange of code and results with others.