

Applied Machine Learning

Parallelization: Batchtools package



Learning goals

- Understand parallelization concepts
- Introduction to batchtools package

PARALLELIZATION



- **Goal:** Minimize computation time by distributing tasks across CPUs/GPUs.
⇒ Speedup is ideally linear but often limited by overhead and dependencies.
- Debugging parallel code is especially hard.
- Coding discipline is even more important to minimize errors and frustration.
- **What can be easily parallelized?**
 - Independent replications
 - Resampling, cross-validation
 - Model averaging
 - Parameter variations in simulations ...
 - "Single program, multiple data"
 - Everything expressible as a loop of independent iterations
(if you can write it with `(1|m|)apply`, you are fine)
- **Many statistical problems are "embarrassingly parallel"**

NAIVE BATCH COMPUTING (NON-CLUSTER)




Workflow on multicore machines:

- Write standalone script(s) to run jobs and save outputs.
- Hard-code parameters or pass via Command-Line Interface (CLI).
- Log in via SSH; run with `R CMD BATCH myscript.R`.
- Use `nohup`, `screen`, or `tmux` to persist after logout.
- Manually start jobs when CPUs free up.
- Check completion and errors by hand.
- Write scripts to merge results.

Drawbacks:

- No resource management, automation, or fair scheduling.
- Poor scalability; hard to debug parallel issues.
- No guarantees for reproducibility (e.g., seeding).

HIGH-PERFORMANCE COMPUTING (HPC) CLUSTERS



www.oxygen-icons.org

- Users access a gateway server (head node).
- Cluster = multiple nodes managed by a scheduler (e.g., SLURM).
- Scheduler assigns jobs to nodes via a queuing system.
- Nodes share a common file system.

MANUAL WORKFLOW ON A HPC CLUSTER



Resource Specification:

- Define CPU count, memory, runtime, partition.
- Set command (e.g., `R CMD BATCH script.R`).

Manual Tasks:

- Submit jobs via CLI or shell scripts.
- Monitor with tools like `squeue`.
- Write aggregation scripts for results.

Typical Workflow:

- Unroll R loops into single-iteration scripts.
- Auto-generate job and script files per task.
- Submit jobs; crawl logs and outputs.

Handling Issues:

- Kill + resubmit on failure.
- Adjust resources on wall-time hits.
- Full rerun for changes in data or params.

BATCHTOOLS OVERVIEW



- R package for structured access to batch systems.
- Built around Map-Reduce: apply algorithms to many problems.
- Full control from R: submit, monitor, kill jobs.
- Persistent state: resume and audit large experiments.
- Convenient debugging and result collection.
- Supports reproducibility across hardware and job schedulers.
- Supports multiple execution backends:
 - **Interactive:** Run jobs directly in the current R session
 - **Multicore:** Parallel execution on local CPU cores
 - **SSH:** Offload jobs to remote machines via SSH
 - **HPC schedulers:** SLURM, Torque/PBS, Load Sharing Facility (LSF), etc.

Project Page: <https://github.com/mllg/batchtools>

Paper: <https://doi.org/10.21105/joss.00135>

CREATING AND CONFIGURING A REGISTRY



Purpose:

- A registry object is used to access and exchange information: file paths, job parameters, and computational events, ...
- Stores all data in a single, portable directory for easy tracking and reproducibility.

Initialization of a new registry:

```
library(batchtools)
reg = makeRegistry(
  file.dir = "registry", # Directory accessible on all nodes
  seed = 1               # Initial seed for reproducibility
)
```

Configure the system:

```
# Set interactive mode and start jobs in external R sessions
reg$cluster.functions = makeClusterFunctionsInteractive(external = TRUE)
```

- Each supported system has its own makeClusterFunctions* function.

Load an existing registry to continue work:

```
loadRegistry("registry")
```

DEFINE JOBS

batchMap:

- Like lapply or mapply
- $(x_1, x_2) \times (y_1, y_2) \rightarrow (f(x_1, y_1), f(x_2, y_2))$
- 10 jobs to calculate $1 + 9, 2 + 8, \dots, 9 + 1$

```
map = function(i, j) i + j  
ids = batchMap(fun = map, i = 1:9, j = 9:1, reg = reg)
```

- Stores function on file system
- Creates jobs as rows in a `data.table`
- Parameters also serialized into the `data.table` for fast access
- All jobs get unique positive integers as IDs
- `reg` = can be omitted in most cases. See `?getDefaultRegistry`.



SUBSET JOBS



Query Job IDs by Status and Parameters

- Use `find*` functions to query job IDs by computational status:
 - `findError` to get job IDs for failed jobs
 - `findDone` to get job IDs for successful jobs
 - `findNotSubmitted` to get job IDs in order resume jobs
 - ...
- Query job IDs by parameters with `findJobs(pars)` (here: `i` and `j`), e.g.,:

```
job = findJobs(i == 2 & j == 8)
job
## Key: <job.id>
##      job.id
##      <int>
## 1:         2
```

- Pass the `data.table` containing the `job.ids` to functions interacting with the batch system, e.g., `submitJobs(ids = job)`

SUBMIT JOBS



- Creates R script files and job description files on the fly
- Resources can be provided as named list

```
# 1 hour maximal execution time, about 2 GB of RAM
res = list(walltime = 60*60, memory = 2000)

# ... and submit
submitJobs(resources = res)
```

- Submits all jobs per default
- Subsets of jobs can be providing as `data.table` or vector

```
submitJobs(ids = 1:5, resources = res)
```

- Collect/reduce results:

```
# get results of each job in a list
reduceResultsList(ids = findDone())

# get result of single job
loadResult(id = 1)
```

SUPERVISE AND DEBUG



- Quick overview of what is going on: `getStatus()`

```
## Status for 9 jobs at 2019-10-10 17:49:48:  
## Submitted      : 9 (100.0%)  
## -- Queued      : 0 (  0.0%)  
## -- Started     : 9 (100.0%)  
## ---- Running   : 0 (  0.0%)  
## ---- Done      : 9 (100.0%)  
## ---- Error     : 0 (  0.0%)  
## ---- Expired   : 0 (  0.0%)
```


- Display log files with a customizable pager (`less`, `vi`, ...):
`showLog(findErrors()[1])`
- You can also `grepLogs(pattern)`
- Found a bug? `killJobs(findRunning())`
- Run a job in the current R session: `testJob(id)`

EXPERIMENTS IN BATCHTOOLS



- **Purpose:** Abstraction for typical statistical tasks.
- **Applying Algorithms to Problems:**
 - Ideal for simulations, benchmark experiments, sensitivity analyses, ...
 - Simplifies workflow with a focus on job definition.
- **Scenarios:**
 - Compare machine learning algorithms on multiple datasets.
 - Compare one/many estimation procedure(s) on simulated data.
 - Compare optimizers on various objective functions.

ABSTRACTION OF COMPUTER EXPERIMENTS



- **Problem Definition:**

- **Static part:** Immutable R objects (e.g., matrices, data frames).
- **Dynamic part:** Arbitrary R functions (e.g., transformations of static objects, data extraction from external sources, data generation functions).

- **Parametrization:** Specify experimental designs for problems and algorithms.

- **Seeding and Reproducibility:**

- Each step is automatically seeded.
- Random seeds are stored in a database for reproducibility.

EXPERIMENT DEFINITION STEPS

- Add problems to registry: `addProblem`
 - Efficient storage: Separation of static (data) and dynamic (instance) problem parts.
- Add algorithms to registry: `addAlgorithm`
 - Problem instance gets passed to algorithm
 - Can be connected with an experimental design (function parameters)
 - Return value will be saved on the file system
- Add experiments to registry: `addExperiments`
 - Experiment: problem instance + algorithm + algorithm parameters
 - Job: Experiment + replication number



A SIMPLE EXAMPLE



```
reg = makeExperimentRegistry("test_reg")
addProblem(name = "p1", data = 1, seed = 1,
  fun = function(data, job) runif(data))
addAlgorithm(name = "a1",
  fun = function(job, data, instance) 2 * instance)
addAlgorithm(name = "a2",
  fun = function(job, data, instance) data + instance)
addExperiments(repls = 2)
submitJobs()
res = reduceResultsDataTable()
getJobPars()[res]
```



```
## Key: <job.id>
##   job.id problem prob.pars algorithm algo.pars  result
##   <int>  <char>    <list>    <char>    <list>  <list>
## 1:      1      p1 <list[0]>      a1 <list[0]> 0.5310
## 2:      2      p1 <list[0]>      a1 <list[0]> 0.3698
## 3:      3      p1 <list[0]>      a2 <list[0]> 1.2655
## 4:      4      p1 <list[0]>      a2 <list[0]> 1.1849
```

SUMMARY



- **Reproducibility:**
 - Every computation is seeded.
 - Seeds are stored in a `data.table`.
- **Extensibility:**
 - Easily add more problems or algorithms.
 - Try different parameters or increase replications at any stage.
- **Portability:** Data, algorithms, results, and job information in a single directory.
- **Exchangeability:** Share your file directory to allow others to extend your study with their data sets and algorithms.
- Simplifies working with batch systems.
- Control batch systems interactively from within R (no shell required).
- Facilitates reproducible research.
- Enables easy exchange of code and results with others.